

GPU Computing

Lab 10

BFS

Parallelo

Struttura di grafo a liste in C++

array **A** offset
dei gradi
sommati

array **E** delle liste
di adiacenza

Restituisce il
grado di **i**

Dice se **j** è nel
vicinato di **i**

```
struct GraphStruct {  
    node nNodes{0}; // num of graph nodes  
    node_sz nEdges{0}; // num of graph edges  
    node_sz* cumulDeps{ nullptr }; // cumsum of node degrees  
    node* neighs{ nullptr }; // list of neighbors for all nodes (edges)  
  
    // return the degree of node i  
    node_sz deg(node i) {  
        return ( cumulDeps[i + 1] - cumulDeps[i] );  
    }  
    // check whether node i is a neighbor of node j  
    bool areNeighbor(node i, node j) {  
        for (unsigned k = 0; k < deg(j); k++)  
            if (neighs[cumulDeps[j]+k] == i)  
                return true;  
        return false;  
    }  
}
```

CUDA BFS

Frontiera:
array F_a

Visitati:
array X_a

Costo:
array C_a

Algorithm 1. CUDA_BFS (Graph $G(V, E)$, Source Vertex S)

- 1: Create vertex array V_a from all vertices and edge Array E_a from all edges in $G(V, E)$,
 - 2: Create frontier array F_a , visited array X_a and cost array C_a of size V .
 - 3: Initialize F_a , X_a to false and C_a to ∞
 - 4: $F_a[S] \leftarrow \text{true}$, $C_a[S] \leftarrow 0$
 - 5: **while** F_a not Empty **do**
 - 6: **for** each vertex V in parallel **do**
 - 7: Invoke CUDA_BFS_KERNEL(V_a, E_a, F_a, X_a, C_a) on the grid.
 - 8: **end for**
 - 9: **end while**
-

BFS Kernel

Frontiera:
array F_a

Visitati:
array X_a

Costo:
array C_a

Algorithm 2. CUDA BFS_KERNEL (V_a, E_a, F_a, X_a, C_a)

```
1:  $tid \leftarrow \text{getThreadID}$ 
2: if  $F_a[tid]$  then
3:    $F_a[tid] \leftarrow \text{false}, X_a[tid] \leftarrow \text{true}$ 
4:   for all neighbors  $nid$  of  $tid$  do
5:     if NOT  $X_a[nid]$  then
6:        $C_a[nid] \leftarrow C_a[tid] + 1$ 
7:        $F_a[nid] \leftarrow \text{true}$ 
8:     end if
9:   end for
10: end if
```

Kernel cudaBFS (hints)

Frontiera:
array F_a

Visitati:
array X_a

Costo:
array C_a

```
__global__ void cudaBFS(GraphStruct *str, bool *Fa, bool *Xa, int *Ca, bool *done, int n) {
    int nodeID = threadIdx.x + blockIdx.x * blockDim.x; // node ID
    if (nodeID > n)
        return;
    if (Fa[nodeID]) {
        *done = false;
        Fa[nodeID] = false;
        Xa[nodeID] = true;
        int deg = str->cumDeps[nodeID + 1] - str->cumDeps[nodeID];
        int start = str->cumDeps[nodeID];
        for (int i = 0; i < deg; i++) {
            int neighID = str->neighs[start + i];
            if ( !Xa[neighID] ) {
                Ca[neighID] = Ca[nodeID] + 1;
                Fa[neighID] = true;
            }
        }
    }
}
```

Merge sort

Parallel sorting in place

Implementazione C

```
/* Main program to sort an array */
int main() {
    . . .

    mergeSort(arr, 0, arr_size - 1);

    . . .
    return 0;
}
```

```
/* l is for left index and r is right index of the
sub-array of arr to be sorted */

void mergeSort(int arr[], int l, int r) {
    if (l < r) {
        // Same as (l+r)/2, but avoids overflow for
        // large l and h
        int m = l + (r - l) / 2;

        // Sort first and second halves
        mergeSort(arr, l, m);
        mergeSort(arr, m + 1, r);


        merge(arr, l, m, r);
    }
}
```

```
void merge(int arr[], int l, int m, int r) {
    int i, j, k;
    int n1 = m - l + 1;
    int n2 = r - m;
    /* create temp arrays */
    int *L = new int[n1];
    int *R = new int[n2];

    /* Copy data to temp arrays L[] and R[] */
    for (i = 0; i < n1; i++)
        L[i] = arr[l + i];
    for (j = 0; j < n2; j++)
        R[j] = arr[m + 1 + j];
    /* Merge the temp arrays back into arr[l..r]*/
    i = 0; j = 0; k = l; // Initial index of subarrays
    while (i < n1 && j < n2) {
        if (L[i] <= R[j]) {
            arr[k] = L[i]; i++;
        } else {
            arr[k] = R[j]; j++;
        }
        k++;
    }
    /* Copy the remaining elements of L[], if there are any */
    while (i < n1) {
        arr[k] = L[i];
        i++; k++;
    }
    /* Copy the remaining elements of R[], if there are any */
    while (j < n2) {
        arr[k] = R[j];
        j++; k++;
    }
}
```


Procedura merge in parallelo

- ✓ Array n-dimensionale ($n = 2^k$)
- ✓ Chunk = $2, 4, 8, \dots, 2^k$



Schema call

Il numero
thread si
dimezza via
via...

Uso di array
alternati per
passaggi
intermedi

. . .

```
for (int chunk = 2; chunk <= N; chunk *= 2) {  
    int nThreads = N / chunk;
```

```
    dim3 block(min(nThreads, BLOCK_SIZE));  
    dim3 grid((nThreads + block.x - 1) / block.x);
```

```
    printf("grid: %d, block: %d, chunk: %d\n", grid.x, block.x, chunk);
```

```
    if (array2sorted)  
        cudaMergeSort<<<grid, block>>>(array, sorted, N, chunk);  
    else  
        cudaMergeSort<<<grid, block>>>(sorted, array, N, chunk);  
    array2sorted = !array2sorted;  
}
```

kernel (1 th per chunk)

chunk = 2,4,8,...

Calcolo estremi
del chunk

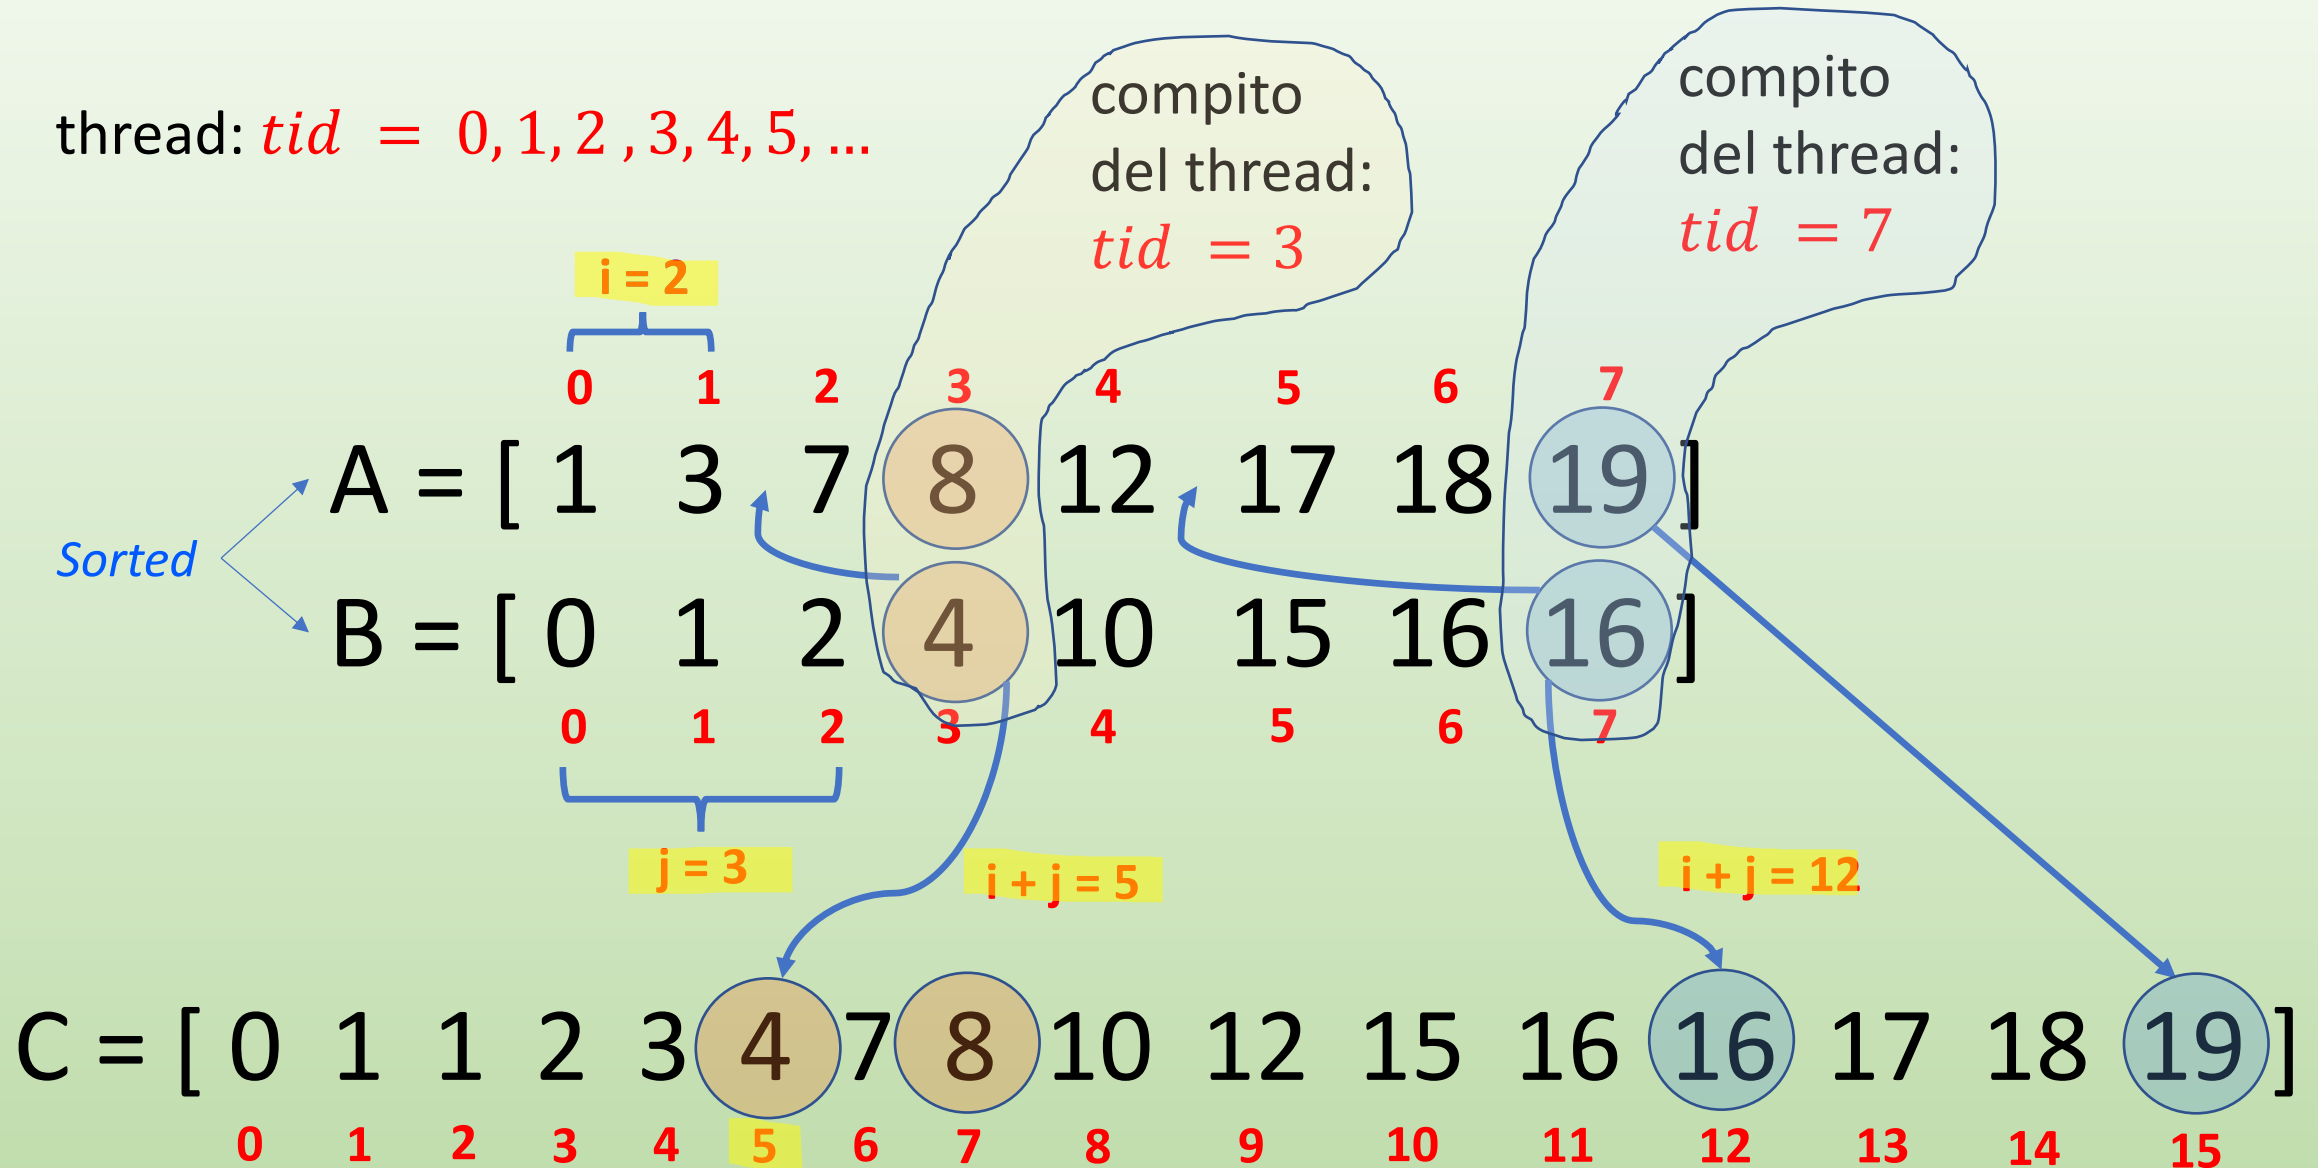
Procedura
merge

Sistema le
'code' eventuali

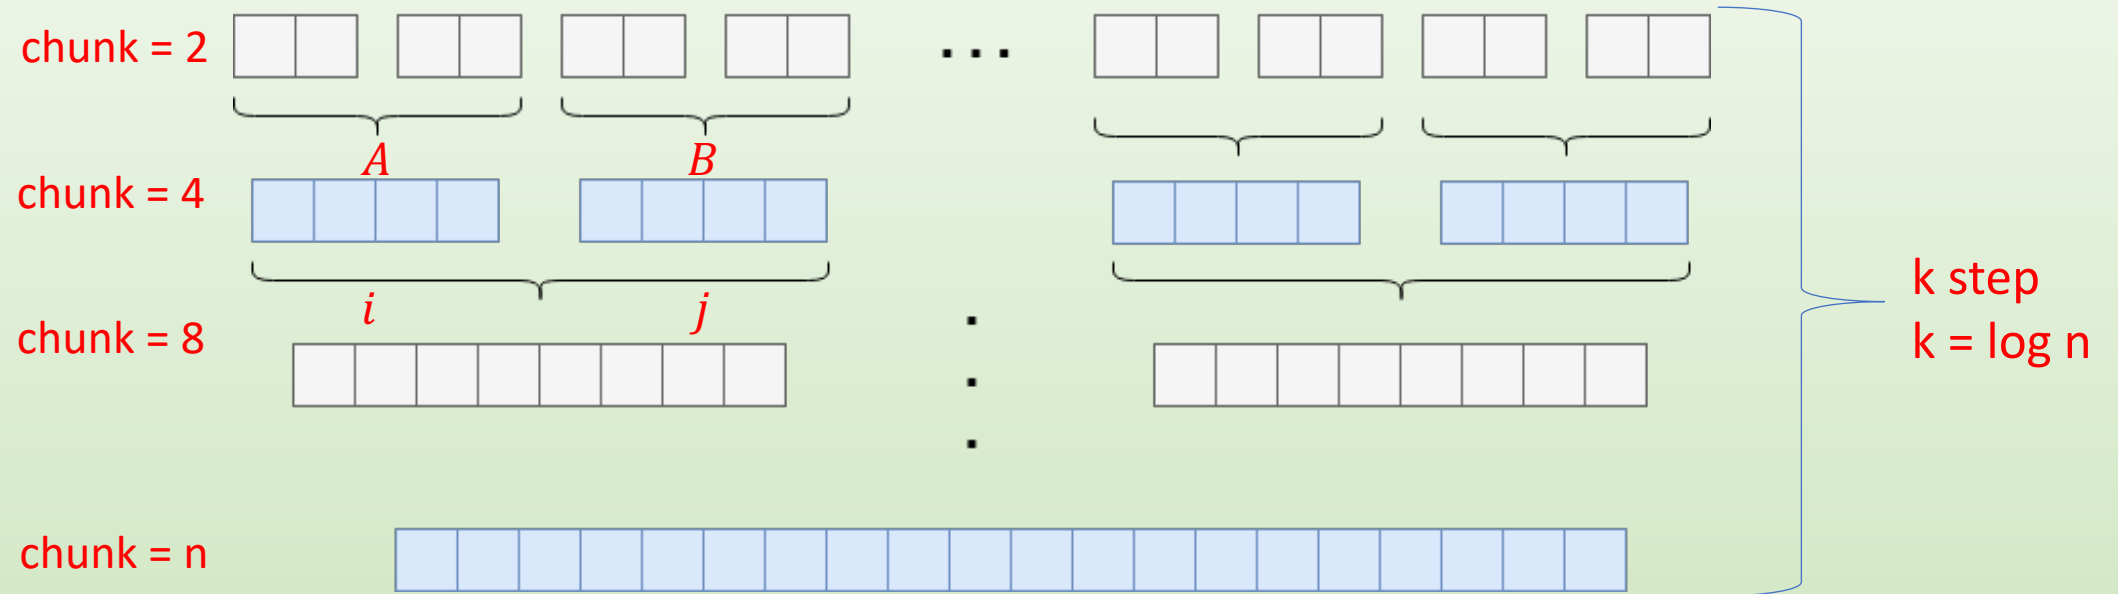
```
...  
__global__ void cudaMergeSort(int *array, int *sorted, int n, int chunk) {  
    int start = chunk * threadIdx.x + blockIdx.x * blockDim.x;  
    if (start > n)  
        return;  
    int mid = min(start + chunk / 2, n);  
    int end = min(start + chunk, n);  
    int i = start, j = end, k = mid;  
  
    //cudaMerge(array, sorted, start, mid, end);  
    while (i < mid && j < end) {  
        if (array[i] <= array[j])  
            sorted[k++] = array[i++];  
        else  
            sorted[k++] = array[j++];  
    }  
    // Copy the remaining elements array[i] if there are any  
    while (i < mid)  
        sorted[k++] = array[i++];  
    // Copy the remaining elements of array[j] if there are any  
    while (j < end)  
        sorted[k++] = array[j++];  
}
```

Parallel binary merge

thread: $tid = 0, 1, 2, 3, 4, 5, \dots$



Procedura merge a più thread



✓ Indice di thread $tid = 0, 1, 2, 3, \dots$

✓ indice elementi $i = tid \% k + (tid - tid \% k) * 2$ (chunk inf A)
 $j = i + k$ (chunk sup B)

Test

```
*** Sorting array size N = 4194304
*** CPU processing... CPU
    elapsed time: 2.69321 (sec)
*** GPU ONE THREAD x chunk processing...
    elapsed time: 0.49219 (sec)
    speedup vs CPU: 5.47
*** GPU MULTI THREAD x chunk processing...
    elapsed time: 0.01057 (sec)
    speedup vs CPU: 254.82
    speedup vs GPU mono: 46.57
```